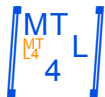




# SimuNova



Cleverer Programmieren  
im Prozessordschungel





- 1 Motivation
  - Software-Entwicklung
  - Ansatz: Generische Programmierung
- 2 MTL4: Intuitiv Programmieren
  - Matrixoperationen
  - Löser
- 3 Zu neuen Ufern: zukunftssichere Anwendungen
  - Portabilität
  - PMTL4: Massiv-parallel
  - CUDA-MTL4: GPUs
  - Performance
- 4 Software-Ökosystem
  - Debugging
  - Nutzer



# Motivation

- Hardware immer diverser
  - Multi-Core, Many-Core,
  - GPU,
  - Cell (RIP),
  - Cluster von all denen
- Verschiedene Speicherstrukturen
  - Hierarchisch auf jeden Fall,
  - Extra Speicher für GPUs,
  - Eventuell verteilt
- Algorithmen und Daten nicht unbedingt homogener



## Motivation

- Hardware immer diverser
  - Multi-Core, Many-Core,
  - GPU,
  - Cell (RIP),
  - Cluster von all denen
- Verschiedene Speicherstrukturen
  - Hierarchisch auf jeden Fall,
  - Extra Speicher für GPUs,
  - Eventuell verteilt
- Algorithmen und Daten nicht unbedingt homogener
- Zukünftige Entwicklung wird nicht einfacher



## Softwarekomplexität reduzieren

- Algorithmen und Datenstrukturen entkoppeln
- Algorithmen und Hardwareaspekte entkoppeln
  - Parallelität abstrahieren
  - Datenlokalität in speziellen Modulen verwalten, nicht in Anwendungen
- Intuitive Schnittstellen
- Wissenschaftler sollen sich nicht mit niederen technischen Details herumschlagen
- Sondern Modelle und Algorithmen entwickeln



## Summierung auf Feldern

```
int sum(int *array, int n)
{
    int s = 0;
    for (int i = 0; i < n; ++i)
        s = s + array[i];
    return s;
}
```

```
double sum(double *array, int n)
{
    double s = 0;
    for (int i = 0; i < n; ++i)
        s = s + array[i];
    return s;
}
```

Das sollte cleverer gehn.



## Templatisierte Summe auf Feldern

```
template <typename T>  
T sum(T *array, int n)  
{  
    T s = 0;  
    for (int i = 0; i < n; ++i)  
        s = s + array[i];  
    return s;  
}
```

Einführung von Templates.



## Summierung verschiedener Container

```
int sum(int *array, int n)
{
    int s = 0;
    for (int i = 0; i < n; ++i)
        s = s + array[i];
    return s;
}

double sum(node *first)
{
    double s = 0;
    while (first) {
        s = s + first->data;
        first = first->next;
    }
    return s;
}
```

Können wir das vereinheitlichen?





## Summierung mit Iteratorbereichen

Einführung des Input-Iterators — jeder Typ mit:

- Inkrementoperator;
- Dereferenzierung und
- Vergleich.

Schreibe Algorithmen für rechtsoffene Intervalle davon:

```
template <typename Iter, typename T>  
T sum(Iter first, Iter last, T s)  
{  
    for ( ; first != last; ++first)  
        s = s + *first;  
    return s;  
}
```



```
int main(int, char**)
{
    mtl::dense2D<double> A, B= {{1, 2}, {3, 4}};
    mtl::matrix<double, morton> C= {{3, 2}, {7, 5}};
    mtl::matrix<float, sparse, col_major> D= {{2, 0}, {0, 3}};

    // Entspricht  $A = B * B$ ; aus Nostalgie; lieber nicht nutzen
    mult(B, B, A);

    A = B * B; // Nimm BLAS (fuer grosse Matrizen)
    A = B * C; // Mit Rekursion + Aufrollen in MTL4

    A += B * D; // Inkrementiere A mit Produkt von B und D
    A -= B * D; // Dekrementiere ...
}
```



```
template <typename LinearOperator, typename SpaceX, typename SpaceB,  
          typename Preconditioner, typename Iteration>  
int cg(const LinearOperator& A, SpaceX& x, const SpaceB& b,  
       const Preconditioner& L, Iteration& iter)  
{  
    typedef typename mtl::Collection<SpaceX>::value_type Scalar;  
    Scalar rho(0), rho_1(0), alpha(0);  
    SpaceX p(resource(x)), q(resource(x)), r(resource(x)), z(resource(x));  
  
    r = b - A*x;  
    while (! iter.finished(r)) {  
        ++iter;  
        z = solve(L, r);  
        rho = dot(r, z);  
  
        if (iter.first())  
            p = z;  
        else  
            p = z + (rho / rho_1) * p;  
  
        q = A * p;  
        alpha = rho / dot(p, q);  
  
        x += alpha * p;  
        r -= alpha * q;  
        rho_1 = rho;  
    }  
    return iter;  
}
```



```
mult(A.data, A.rows, x.data, ...)  
sub(y.data, y.rows, ...)
```

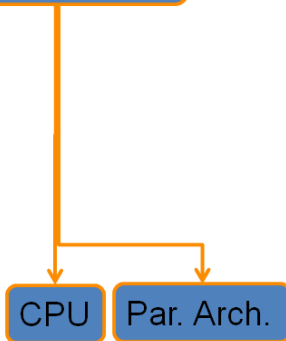
Scientific app.

CPU



```
mult(A.data, A.rows, x.data, ...)  
sub(y.data, y.rows, ...)  
MPI_Send(A.ldata, lsize, ...)  
distr_mult(A.ldata, y.rows, ...)
```

Scientific app.  
with Parallelism

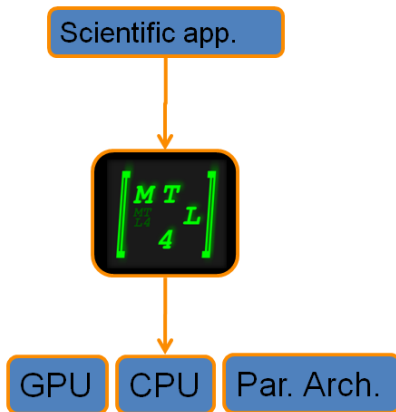




```
mult(A.data, A.rows, x.data, ...)  
sub(y.data, y.rows, ...)  
MPI_Send(A.ldata, lsize, ...)  
distr_mult(A.ldata, y.rows, ...)  
CUDA_copy(&A.data, ...)  
mult_kernel(&A.ldata, ...)
```

Scientific app.  
with Parallelism  
& some GPU support







$r = b - A^*x;$   
 $q = P^*r;$

Scientific app.



GPU

CPU

Par. Arch.





$r = b - A^*x;$   
 $q = P^*r;$

Scientific app.





$r = b - A^*x;$   
 $q = P^*r;$

Scientific app.



Meta-tuning[GS10]  
Platform conceptualization  
Property-aware programming[GL06]



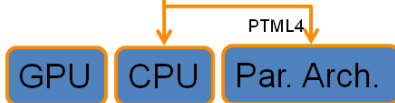


$r = b - A^*x;$   
 $q = P^*r;$

Scientific app.



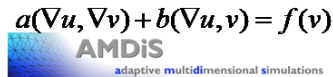
Meta-tuning[GS10]  
Platform conceptualization  
Property-aware programming[GL06]





$r = b - A^*x;$   
 $q = P^*r;$

Scientific app.



Meta-tuning[GS10]  
Platform conceptualization  
Property-aware programming[GL06]



GTML4

PTML4

GPU

CPU

Par. Arch.

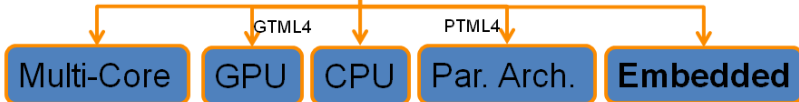


$r = b - A^*x;$   
 $q = P^*r;$

Scientific app.



Meta-tuning[GS10]  
Platform conceptualization  
Property-aware programming[GL06]



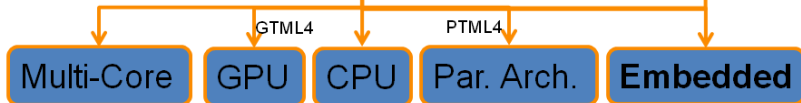


$r = b - A^*x;$   
 $q = P^*r;$

Scientific app.



Meta-tuning[GS10]  
Platform conceptualization  
Property-aware programming[GL06]





```
int main(int argc, char* argv[])
{
    using namespace mtl;
    par::environment env(argc, argv);

    matrix<float, sparse> A, B(12, 12), C;
    laplacian_setup(A, 3, 4);
    B = 2.0;

    C = 2.0 * A - B / 2.0;
    C += A / 4.0 + conj(B);
    C += A * B;

    return 0;
}
```



```
int main(int argc, char* argv[])
{
    using namespace mtl;
    using vector_type = mtl::vector<float>;
    using matrix_type = matrix<float>;

    par::environment env(argc, argv);
    vector_type b(12, 1.0);
    matrix_type A;
    laplacian_setup(A, 4, 3);

    auto migration = parmetis_migration(A);
    vector_type b2(b, migration), x2(resource(b2));
    matrix_type A2(A, migration);
    // Solve A2 * x2 = b2

    vector_type x(x2, reverse(migration));
}
```





## Anwendung Kristallwachstum

- Parallel auf zwei Cloud-Installationen in China,
- Parametrisierung mühsamer als Programmierung,
- <https://www.youtube.com/watch?v=8v6snbjC38Q>
- <https://www.youtube.com/watch?v=xv9J0UCieNQ>



```
--global__ void
add(int *a, int *b, int *c ) {
    int tid = blockIdx.x; // handle the data at this index
    if (tid < N)
        c[tid] = a[tid] + b[tid];
}

#define N 10
int main(void) {
    int a[N], b[N], c[N];
    int *dev_a, *dev_b, *dev_c;
    HANDLE_ERROR( cudaMalloc( (void**)&dev_a, N * sizeof(int) ) );
    HANDLE_ERROR( cudaMalloc( (void**)&dev_b, N * sizeof(int) ) );
    HANDLE_ERROR( cudaMalloc( (void**)&dev_c, N * sizeof(int) ) );

    for (int i=0; i<N; i++) {
        a[i] = -i; b[i] = i * i;
    }
    HANDLE_ERROR( cudaMemcpy( dev_a, a, N * sizeof(int), cudaMemcpyHostToDevice ) );
    HANDLE_ERROR( cudaMemcpy( dev_b, b, N * sizeof(int), cudaMemcpyHostToDevice ) );
    HANDLE_ERROR( cudaMemcpy( dev_c, c, N * sizeof(int), cudaMemcpyHostToDevice ) );

    add<< <N,1 >>>( dev_a, dev_b, dev_c );

    HANDLE_ERROR( cudaMemcpy( c, dev_c, N * sizeof(int), cudaMemcpyDeviceToHost ) );
    for(int i=0; i<N; i++)
        printf(" %d + %d = %d\n", a[i], b[i], c[i] );

    cudaFree( dev_a ); cudaFree( dev_b ); cudaFree( dev_c );
    return 0;
}
```



## Abstrakte (CUDA-) Programmierung

```
#include <iostream>
#include <boost/numeric/mtl/mtl.hpp>

int main(int argc, char** argv)
{
    const int N= 10;
    mtl::dense_vector<int> a(N), b(N), c;
    for (int i: irange(0, N))
        a[i]= -i, b[i]= i * i;
    c= a + b;
    std::cout << a << " + " << b << " = " << c << "\n";
    return 0;
}
```

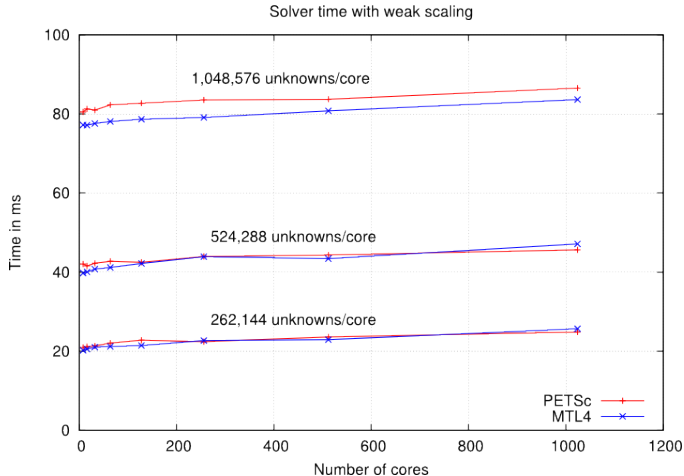


Abbildung : Vergleich parallele CG-Performance mit PETSc

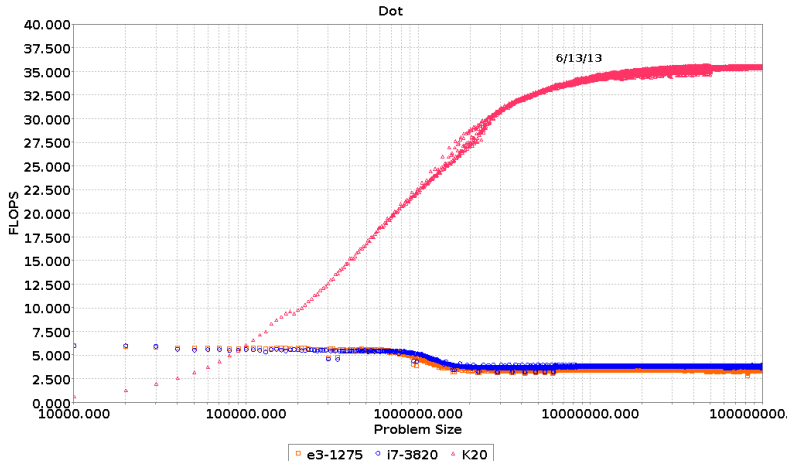


Abbildung : Skalarprodukt auf CPU und GPU

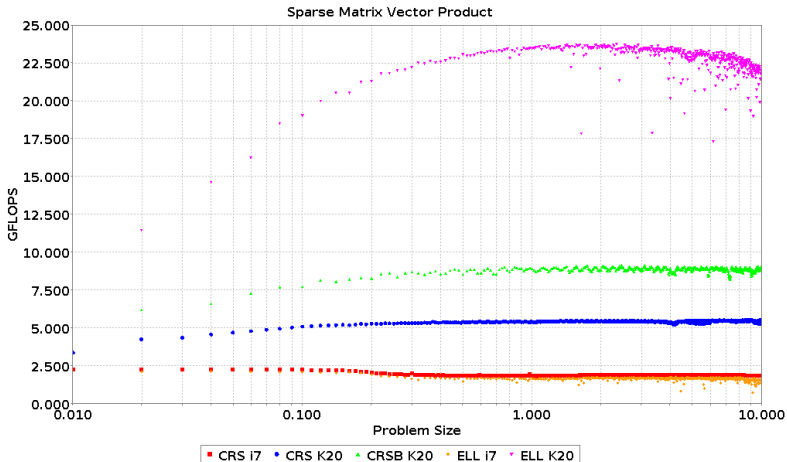


Abbildung : Matrix-Vektor-Produkt

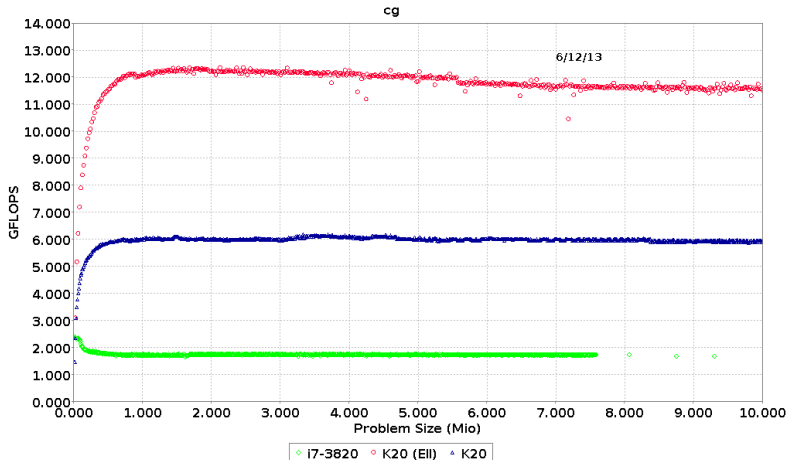
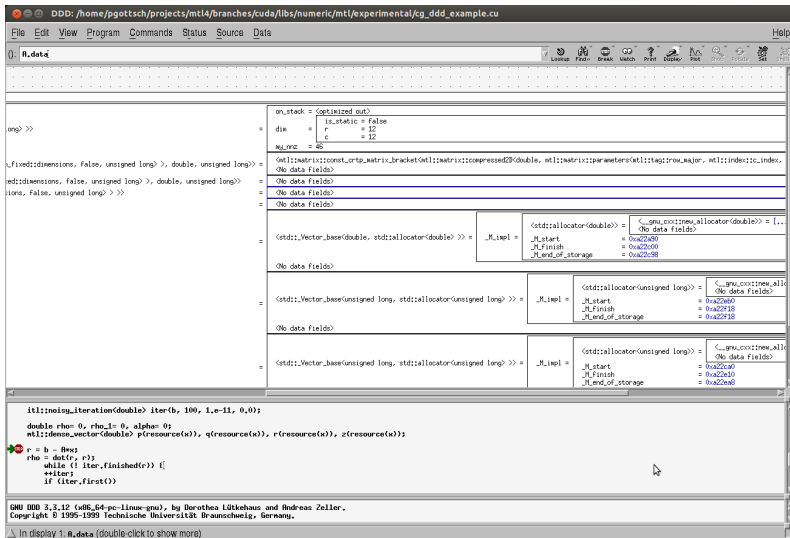
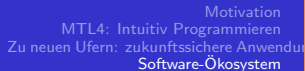


Abbildung : Konjugierte Gradienten







DDD: /home/pgottsch/projects/mtl4/branches/cuda/libs/numeric/mtl/experimental/cg\_ddd\_example.cu

File Edit View Program Commands Status Source Data Help

1: /cuda/libs/numeric/mtl/experimental/cg\_ddd\_example.cu:38

1: A

4	-1	0	0	-1	0	0	0	0	0	0	0
-1	4	-1	0	0	-1	0	0	0	0	0	0
0	-1	4	-1	0	0	-1	0	0	0	0	0
0	0	-1	4	0	0	0	-1	0	0	0	0
-1	0	0	0	4	-1	0	0	-1	0	0	0
0	-1	0	0	-1	4	-1	0	0	-1	0	0
0	0	-1	0	0	-1	4	-1	0	0	-1	0
0	0	0	-1	0	0	-1	4	0	0	-1	0
0	0	0	0	-1	0	0	0	4	-1	0	0
0	0	0	0	0	-1	0	0	-1	4	-1	0
0	0	0	0	0	0	-1	0	0	-1	4	-1
0	0	0	0	0	0	0	-1	0	0	-1	4

2: b

device
2
1
1
2
1
0
0
1
2
1
1
1
2

3: x

device
0
0
0
0
0
0
0
0
0
0
0
0
0
0

5: r

host
2
1
1
2
1
0
0
0
1
2
1
1
1
2

4: p

host
6.9533485240516674e-310
6.9533485240516674e-310
0
0
1.3580773062177743e-312
2.1219957904712067e-314
0
0
1.379297264127427e-312
2.1219957904712067e-314
0
0
0
2.6529260140087942e+180

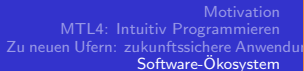
6: q

host
6.9533485240516674e-310
6.9533485240516674e-310
0
0
9.6895420172743119e-71
1.5955166634036102e+184
1.1883710132235731e-81
0
1.9410091417419864e-80
4.546040660768358e+150
1.811525845712515e+272
1.033315611527298e-47
2.6529260140087942e+180

```
double rho(0), rho_1(0), alpha(0), alpha_1(0);
mtl::dense_vector<double> p(resource(x)), q(resource(x)), r(resource(x)), z(resource(x));

r = b - A*x;
rho = dot(r, r);
while (! iter.finished(r)) {
  ++iter;
  if (iter.first())
    p = r;
  else
    p = r + (rho / rho_1) * p;
  q = A * p; alpha_1 = dot(p, q);
  alpha= rho / alpha_1;
  x += alpha * p;
  rho_1 = rho;
}
```

(gdb) next  
(gdb) next  
iteration 0: resid 4.69042  
(gdb) next  
(gdb) next  
(gdb)



```
(gdb) next
iteration 0: resid 4.69042
(gdb) next
(gdb) next
(gdb) next
(gdb) next
```



DDD: /home/pgottsch/projects/mtl4/branches/cuda/libs/numeric/mtl/experimental/cg\_ddd\_example.cu

File Edit View Program Commands Status Source Data Help

1: A

4	-1	0	0	-1	0	0	0	0	0	0	0
-1	4	-1	0	0	-1	0	0	0	0	0	0
0	-1	4	-1	0	0	-1	0	0	0	0	0
0	0	-1	4	0	0	0	-1	0	0	0	0
-1	0	0	0	4	-1	0	0	-1	0	0	0
0	-1	0	0	-1	4	-1	0	0	-1	0	0
0	0	-1	0	0	-1	4	-1	0	0	-1	0
0	0	0	-1	0	0	-1	4	0	0	0	-1
0	0	0	0	-1	0	0	0	4	-1	0	0
0	0	0	0	0	-1	0	0	-1	4	-1	0
0	0	0	0	0	0	-1	0	0	-1	4	-1
0	0	0	0	0	0	0	-1	0	0	-1	4

2: b

device
2
0
1
1
0
2
0
1
0
0
0
1
2
0
1
1
2

3: x

device
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0

5: r

host
2
0
1
1
0
2
0
1
0
0
0
1
2
0
1
1
2

4: p

device
2
1
1
1
2
1
0
0
0
1
2
2
1
1
1
2

6: q

device
6
1
1
6
0
-3
-3
0
6
1
1
1
6

```
double rho(0), rho_1(0), alpha(0), alpha_1(0);
mtl::dense_vector<double> p(resource(x)), q(resource(x)), r(resource(x)), z(resource(x));

r = b - A*x;
rho = dot(r, r);
while (! iter.finished(r)) {
  ++iter;
  if (iter.first())
    p = r;
  else
    p = r + (rho / rho_1) * p;

  q = A * p; alpha_1 = dot(p, q);
  alpha = rho / alpha_1;

  x += alpha * p;
  rho_1 = rho;
}
```

iteration 0: resid 4.69042  
(gdb) next  
(gdb) next  
(gdb) next  
(gdb) next  
(gdb)



DDD: /home/pgottsch/projects/mtl4/branches/cuda/libs/numeric/mtl/experimental/cg\_ddd\_example.cu

File Edit View Program Commands Status Source Data Help

1: /cuda/libs/numeric/mtl/experimental/cg\_ddd\_example.cu:47

1: A

4	-1	0	0	-1	0	0	0	0	0	0	0
-1	4	-1	0	0	-1	0	0	0	0	0	0
0	-1	4	-1	0	0	-1	0	0	0	0	0
0	0	-1	4	0	0	0	-1	0	0	0	0
-1	0	0	0	4	-1	0	0	-1	0	0	0
0	-1	0	0	-1	4	-1	0	0	-1	0	0
0	0	-1	0	0	-1	4	-1	0	0	-1	0
0	0	0	-1	0	0	-1	4	0	0	0	-1
0	0	0	0	-1	0	0	0	4	-1	0	0
0	0	0	0	0	-1	0	0	-1	4	-1	0
0	0	0	0	0	0	-1	0	0	-1	4	-1
0	0	0	0	0	0	0	-1	0	0	-1	4

2: b

device
2
0,84615384615384615
0,42307692307692307
1
0,42307692307692307
0,84615384615384615
0,42307692307692307
0
0
1
0
2
0,84615384615384615
0,42307692307692307
0,42307692307692307
0,84615384615384615

3: x

device
2
0,84615384615384615
0,42307692307692307
1
0,42307692307692307
0,84615384615384615
0,42307692307692307
0
0
1
0
2
0,84615384615384615
0,42307692307692307
0,42307692307692307
0,84615384615384615

5: r

device
2
-0,53846153846153832
0,57692307692307687
0,57692307692307687
-0,53846153846153832
1
1,2692307692307692
1,2692307692307692
1
-0,53846153846153832
0,57692307692307687
0,57692307692307687
-0,53846153846153832

4: p

device
2
1
1
1
2
0
0
1
2
1
1
2

6: q

device
6
1
1
6
0
-3
-3
0
6
1
1
6

```
rho = dot(r, r);
while (! iter.finished(r)) {
  ++iter;
  if (iter.first())
    p = r;
  else
    p = r + (rho / rho_1) * p;

  q = A * p; alpha_1 = dot(p, q);
  alpha = rho / alpha_1;

  x += alpha * p;
  rho_1 = rho;
  r -= alpha * q; rho = dot(r, r);
}
return 0;
}
```

(gdb) next  
(gdb) next  
(gdb) next  
(gdb) next  
(gdb) next  
(gdb)



- FEniCS: freie Software zum automatischen Lösen differentieller Gleichungen.
  - University of Chicago, Argonne National Laboratory, Delft University of Technology, Royal Institute of Technology KTH, Simula Research Laboratory, Texas Tech University und University of Cambridge.
- Bombardier: Entwicklung neuer Flugzeuge.
- Institut Français du Pétrole: Mehrphasensimulation unterirdischer CO<sub>2</sub>-Speicherung mittels finiter Elemente.
- Serbisches Institut für Grundwasserressourcen: dichte-getriebene Grundwasserströmungen mit stark diskontinuierlicher und anisotroper Permeabilität.